

## Day2

### 1- Creating s3 for state file by terraform

```
aws_s3_bucket.terraform_state: Creating ...
aws_s3_bucket.terraform_state: Creation complete after 8s [id=dev-bucket-state-lab2.0]

Apply complete! Resources: 1 added, 0 changed, 0 destroyed.
```

The screenshot shows the AWS S3 console interface. At the top, there are two tabs: "General purpose buckets" (which is selected) and "Directory buckets". Below the tabs, it says "General purpose buckets (1) [Info](#)". There are four buttons: "Copy ARN", "Empty", "Delete", and "Create bucket". A note below the buttons states: "Buckets are containers for data stored in S3.". A search bar is present with the placeholder "Find buckets by name". To the right of the search bar are navigation arrows and a gear icon. A table below lists the bucket details:

Name	AWS Region	Creation date
<a href="#">dev-bucket-state-lab2.0</a>	US East (N. Virginia) us-east-1	November 2, 2025, 20:25:02 (UTC+02:00)

### 2- organize Terraform configurations into logical files (versions.tf, provider.tf, backend.tf, variables.tf, and main.tf).

```
[fayyad@rocky day2]$ ls
backend.tf  main.tf  outputs.tf  provider.tf  variables.tf  versions.tf
```

### 3- use the terraform block to define required providers and version constraints.

```
terraform {
  required_version = "≥ 1.6.0"

  required_providers {
    aws = {
      source  = "hashicorp/aws"
      version = "≈ 5.0"
    }
  }
}
```

4- use variables for EC2 instance configuration.

```
variable "aws_region" {
  description = "The AWS region to deploy the resources"
  type        = string
  default     = "us-east-1"
}

variable "instance_type" {
  description = "The EC2 instance type"
  type        = string
  default     = "t2.micro"
}

variable "instance_name" {
  description = "The Name tag for the EC2 instance"
  type        = string
  default     = "Terraform-Lab-EC2"
}

variable "instance_count" {
  description = "(Bonus) Number of EC2 instances to deploy"
  type        = number
  default     = 1
}
```

5- use a data source to dynamically retrieve the latest Amazon Linux 2 AMI ID.

```
# Data source to get the latest Amazon Linux 2 AMI

data "aws_ami" "amazon_linux_2" {
  most_recent = true
  owners      = ["amazon"]

  filter {
    name  = "name"
    values = ["amzn2-ami-hvm-*x86_64-gp2"]
  }

  filter {
    name  = "virtualization-type"
    values = ["hvm"]
  }
}
```

6- configure an S3 remote backend with the use\_lockfile flag for state management and locking.

```
terraform {
  backend "s3" {
    bucket      = "dev-bucket-state-lab2.0"
    key         = "ec2-deployment/terraform.tfstate"
    region      = "us-east-1"
    use_lockfile = "true"
  }
}
```

7- output important information such as instance public IP addresses.

```
output "instance_public_ips" {
  description = "Public IP addresses of the deployed EC2 instance(s)"
  value       = aws_instance.lab_instance.*.public_ip # Uses a splat expression to get all IPs
}

output "instance_ids" {
  description = "IDs of the deployed EC2 instance(s)"
  value       = aws_instance.lab_instance.*.id
}
```